

M2U3 – Individual Assignment: RAG Mini Project (n8n)

 [github](#) [osama-ata/MAICEN2511](#)

1) Goal


Build a simple Retrieval-Augmented Generation (RAG) workflow in **n8n** that answers questions **only using my own dataset**. The workflow loads structured project data, filters it into relevant subsets (retrieved context), and provides that context to an LLM to generate a constrained answer.

2) Submission Artifacts (Deliverables)

- **Workflow export (JSON):** M2U3/RAG_Project_Workflow.json
- **Screenshots (3–6 total):** M2U3/screenshots/

Screenshot per major section of the workflow:

- 01_workflow.png
- 02_getting_data.png
- 03_filtering_data.png
- 04_merging_datasets.png
- 05_merging_prompt_context.png
- 06_llm_call.png



Overall Workflow

```

[{"site_id": "S001", "site_name": "Downtown Plaza", "estimated_materials_cost": 45000, "actual_materials_cost": 47500, "estimated_labor_cost": 32000, "actual_labor_cost": 34500, "estimated_duration_days": 90, "actual_duration_days": 95}, {"site_id": "S002", "site_name": "Harbor View Complex", "estimated_materials_cost": 78000, "actual_materials_cost": 82000, "estimated_labor_cost": 56000, "actual_labor_cost": 59000, "estimated_duration_days": 120, "actual_duration_days": 125}, {"site_id": "S003", "site_name": "Westside Mall", "estimated_materials_cost": 120000, "actual_materials_cost": 115000, "estimated_labor_cost": 85000, "actual_labor_cost": 88000}
]

```

Getting Data

INPUT

Parameters

Conditions

```

fx {{ $json.estimated_materials_cost }} # is less than
fx {{ $json.actual_materials_cost }}

```

Add condition

Options

No properties

Add option

OUTPUT

True Branch (102 items) False Branch (3 items)

```

[{"site_id": "S001", "site_name": "Downtown Plaza", "estimated_materials_cost": 45000, "actual_materials_cost": 47500, "estimated_labor_cost": 32000, "actual_labor_cost": 34500, "estimated_duration_days": 90, "actual_duration_days": 95}, {"site_id": "S002", "site_name": "Harbor View Complex", "estimated_materials_cost": 78000, "actual_materials_cost": 82000, "estimated_labor_cost": 56000, "actual_labor_cost": 59000, "estimated_duration_days": 120, "actual_duration_days": 125}, {"site_id": "S003", "site_name": "Westside Mall", "estimated_materials_cost": 120000, "actual_materials_cost": 115000, "estimated_labor_cost": 85000, "actual_labor_cost": 82000, "estimated_duration_days": 150, "actual_duration_days": 145}, {"site_id": "S004", "site_name": "Riverside Tower", "estimated_materials_cost": 95000, "actual_materials_cost": 102000, "estimated_labor_cost": 68000, "actual_labor_cost": 73000, "estimated_duration_days": 110, "actual_duration_days": 118}], [{"site_id": "S001", "site_name": "Downtown Plaza", "estimated_materials_cost": 45000, "actual_materials_cost": 47500, "estimated_labor_cost": 32000, "actual_labor_cost": 34500, "estimated_duration_days": 90, "actual_duration_days": 95}, {"site_id": "S002", "site_name": "Harbor View Complex", "estimated_materials_cost": 78000, "actual_materials_cost": 82000, "estimated_labor_cost": 56000, "actual_labor_cost": 59000, "estimated_duration_days": 120, "actual_duration_days": 125}, {"site_id": "S004", "site_name": "Riverside Tower", "estimated_materials_cost": 95000, "actual_materials_cost": 102000, "estimated_labor_cost": 68000, "actual_labor_cost": 73000, "estimated_duration_days": 110, "actual_duration_days": 118}]

```

Page Size 25

Filtering Data

INPUT

Parameters

Mode

Append

OUTPUT

This node will output an empty item if nothing would normally be returned

```

3 items
[{"Sites with Material Cost Overrun": [...], "Sites with Labor Cost Overrun": [...], "Delayed Project": [...]}]

```

Page Size 25

Merging Datasets

The screenshot shows the n8n interface with the AI Agent node selected. The 'Parameters' tab is active, showing a tip about prompt context merging, the source for prompt user message (Connected Chat Trigger Node), and the prompt (User Message) which includes a JSON template and fallback text. The 'Options' tab shows system messages for the AI agent. The right side shows the output section with 'No output data'.

Merging Prompt Context

The screenshot shows the n8n interface with the Logs - RAG Project - n8n interface. It displays a chat log where the user asks "Which projects are delayed?" and the AI Agent responds with a list of delayed projects. The AI Agent node is expanded to show its configuration, including the OpenAI Chat Model and its parameters like messages, options (openai_api_key), configuration (baseURL), and model_kwargs.

LLM Call

3) Data Used

- **Dataset:** construction project/site cost + duration data (JSON derived from CSV)
 - **Primary columns:**
 - site_id , site_name
 - estimated_materials_cost , actual_materials_cost
 - estimated_labor_cost , actual_labor_cost
 - estimated_duration_days , actual_duration_days
 - **Source file referenced by workflow:** M2U3/project_data.json (fetched via HTTP in the workflow)
-

4) Workflow Summary (What happens)

1. Load the dataset (JSON).
2. “Retrieve” relevant subsets by filtering:
 - sites where **actual materials cost > estimated**
 - sites where **actual labor cost > estimated**
 - sites where **actual duration > estimated**
3. Aggregate filtered results into a single **context** object.
4. Provide that context to the LLM so answers are grounded in retrieved data.

This design supports questions such as:

- “Which sites have material cost overruns?”
 - “Which sites are delayed?”
 - “List sites with labor overruns.”
-

5) One Sentence Per Node (as required)

- **When chat message received:** Receives the user’s question (chat input) to trigger the workflow.
- **get_projects_data:** Loads the project dataset (JSON) into the workflow.

- **materials_cost_overrun (IF)**: Filters rows where `actual_materials_cost` exceeds `estimated_materials_cost`.
 - **labor_cost_overrun (IF)**: Filters rows where `actual_labor_cost` exceeds `estimated_labor_cost`.
 - **delayed_projects (IF)**: Filters rows where `actual_duration_days` exceeds `estimated_duration_days`.
 - **Aggregate_materials_cost_overrun**: Collects the `site_name` values for material-overrun sites into one list.
 - **Aggregate_labor_cost_overrun**: Collects the `site_name` values for labor-overrun sites into one list.
 - **aggregate_delayed_projects**: Collects the `site_name` values for delayed sites into one list.
 - **merge_cost**: Merges the two cost-overrun aggregations into a single stream.
 - **merge_cost_time**: Merges cost results with delay results to form a combined dataset.
 - **context (Aggregate)**: Aggregates merged results into a single `context` payload for the LLM.
 - **OpenAI Chat Model1**: Provides the chat model used by the agent for generation.
 - **AI Agent**: Generates the final answer using only the provided `context`.
-

6) RAG Prompt Used

Answer only using the data below.

DATA:

`{{context}}`

QUESTION:

`{{query}}`

If the answer is not in the data, reply exactly: "No information available."

Note: In this workflow, the “retrieved context” is the filtered + aggregated lists produced upstream (overruns/delays), which constrains what the model can truthfully answer.

7) How to Run (n8n)

1. Open n8n.
 2. Import the workflow JSON: M2U3/RAG_Project_Workflow.json
 3. Configure credentials for the OpenAI node (or replace with your preferred provider).
 4. Execute the workflow (or use the chat trigger) and ask questions like:
 - “Which sites are delayed?”
 - “Which sites have labor cost overruns?”
-

8) Short Reflection (\leq 500 words)

What the data is:

The dataset represents construction sites with estimated vs. actual costs (materials and labor) and estimated vs. actual durations. Each row is a site, which makes it suitable for retrieval-by-filtering.

What questions people will ask:

Typical stakeholder questions are list-based and exception-focused: which sites are over budget (materials/labor), which are delayed, and which sites require attention due to negative variance.

What the AI should reply with:

Short, factual answers grounded in the retrieved results (e.g., a list of site names that match the condition). If a question cannot be answered from the retrieved context, the assistant must respond exactly with: **“No information available.”**

Why clean data matters:

RAG depends on matching and filtering. If column names, types, or units are inconsistent (e.g., costs stored as strings, missing values, mixed currencies), retrieval will either miss relevant rows or include incorrect ones—both of which degrade the quality of the final answer.

Why filtering is “search before generate”:

The workflow first narrows the dataset to only the rows that meet the query-relevant conditions (overrun or delay). This step plays the role of retrieval: it reduces the information

space before the model generates any natural language output.

How RAG reduces hallucination:

By supplying a bounded context (filtered lists) and instructing the model to only use that context, the model has less room to invent unsupported details.

What was easy vs. confusing:

Using n8n's required extra learning to understand how to chain nodes and pass data between them. However, the visual interface made it easier to conceptualize the workflow steps compared to coding from scratch.